



Chap. 4

ARM Organization and Implementation

4.1 3-stage pipeline ARM organization

■ The 3-stage pipeline

- ◆ ARM processors up to the ARM7 employ a simple 3-stage pipeline with the following pipeline stages: (Fig. 4.1)
 - *Fetch*: The instruction is fetched from memory and placed in the instruction pipeline
 - *Decode*: The instruction is decoded and the datapath control signals prepared for the next cycle
 - *Execute*: The register bank is read, an operand shifted, the ALU result generated and written back into a destination register
- ◆ An individual instruction takes three clock cycles to complete, so it has a three-cycle *latency*, but the *throughput* is one instruction per cycle (Fig. 4.2)





Fig. 4.2 ARM single-cycle instruction pipeline operation

Pipeline

- 3 stages (ARM7) and 5 stages (ARM9TDMI)



PC
*Load an instruction
from memory*

PC-4

PC-8

*access memory
if needed*

*write result
to register*




Fig. 4.1 ARM organization

4.1 3-stage pipeline ARM organization

■ The 3-stage pipeline (Cont.)

- ◆ When a multi-cycle instruction is executed the flow is less regular




Fig. 4.3 ARM multi-cycle instruction pipeline operation

4.1 3-stage pipeline ARM organization

■ The 3-stage pipeline (Cont.)

- ◆ All instructions occupy the datapath for one or more adjacent cycles
- ◆ For each cycle that an instruction occupies the datapath, it occupies the decode logic in the immediately preceding cycle
- ◆ During the first datapath cycle each instruction issues a fetch for the next instruction but one
- ◆ Branch instructions flush and refill the instruction pipeline

■ Due to the pipelined execution model used on the ARM is that the program counter, which is visible to the user as r15, must run ahead of the current instruction.

- ◆ PC must point eight bytes ahead of the current executed instruction.
- ◆ In Chapter5:
 - Specifying the PC in an STM should be avoided

5.2 Exceptions

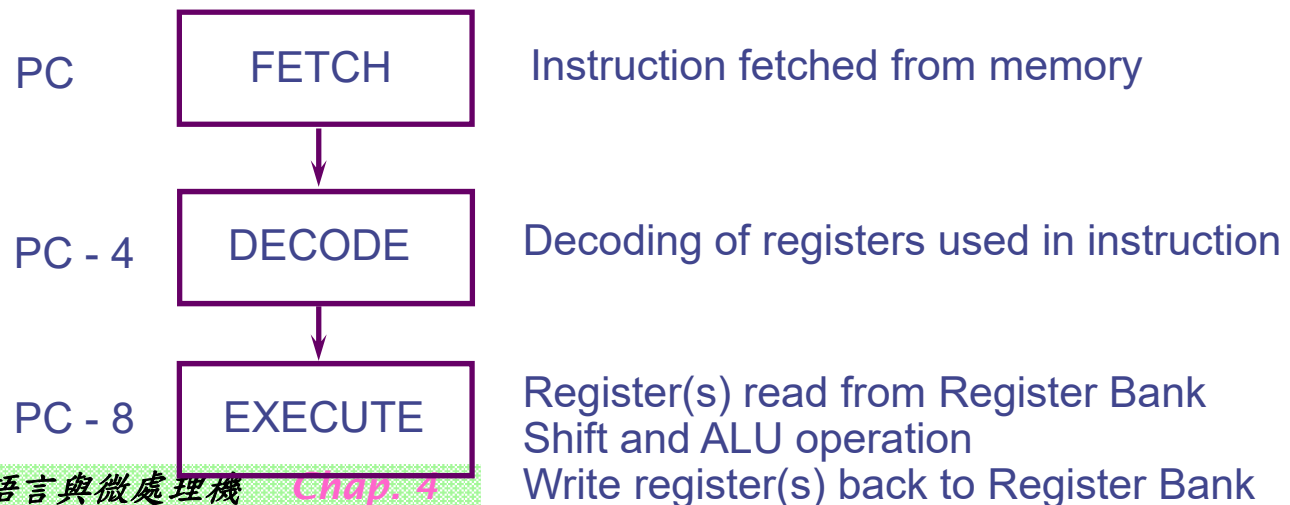
■ Exception return

- ◆ Any **modified user registers** must be **restored** from the handler's stack
- ◆ The **CPSR** must be **restored** from the appropriate SPSR
- ◆ The **PC** must be changed **back** to the relevant instruction address in the user instruction stream
- ◆ The last two steps cannot be carried out independently.
- ◆ The case where the return address is in r14 (S: destination is pc):
 - To return from a SWI or undefined instruction trap use
MOVS pc, r14
 - To return from an IRQ, FIQ or prefetch abort use
SUBS pc, r14, #4
 - To return from a data abort to retry the data access use
SUBS pc, r14, #8
- ◆ The return address out onto a stack:
 - LDMFD r13!, {r0 – r3, pc}^ ; restore and return CPSR is also restore

The Instruction Pipeline

Exception	Address	Use
Reset	-	Ir is not defined
Data abort	lr-8	Points to the instructions that caused data abort
FIQ	lr-4	Return to next instruction
IRQ	lr-4	Return to next instruction
Prefetch Abort	lr-4	Points to the instruction that cause prefetch abort
SWI	lr	Return to next instruction after SWI
Undefined Instruction	lr	Return to next instruction after undefined instruction

ARM



4.2 5-stage pipeline ARM organization

■ 3-stage (cost-effective) \Rightarrow 5-stage (higher performance)

- ◆ The time, T_{prog} , required to execute a given program is given by:

$$T_{prog} = N_{inst} \times CPI / f_{clk}$$

- N_{inst} : number of ARM instructions executed in the course of the program
 - CPI : the average number of clock cycles per instruction
 - f_{clk} : the processor's clock frequency
-
- ◆ There are two ways to increase performance:
 - Increase the clock rate, f_{clk}
 - Logic in each pipeline stage to be simplified \Rightarrow the number of pipeline stage increased
 - Reduce the average number of clock cycles per instruction, CPI
 - Instructions which occupy more than one pipeline slot in a 3-stage pipeline ARM are re-implemented to occupy fewer slots
 - Pipeline stalls caused by dependencies between instruction are reduced

4.2 5-stage pipeline ARM organization

■ Memory bottleneck

- ◆ Any stored-program computer with a single instruction and data memory will have its performance limited by the available **memory bandwidth**
- ◆ To get a significantly better **CPI** the memory system must deliver more than one value in each clock cycle either by delivering more than 32 bits per cycle from a single memory or by having **separate memories for instruction and data accesses**
- ◆ Higher performance ARM cores employ a **5-stage pipeline** and have **separate instruction and data memories**

4.2 5-stage pipeline ARM organization

■ The 5-stage pipeline

◆ *Fetch*

- the **instruction** is fetched from memory and placed in the **instruction pipeline**

◆ *Decode*

- the instruction is **decoded** and **register operands** read from the register file

◆ *Execute*

- an operand is shifted and the **ALU** result generated

◆ *Buffer/data*

- **data memory** is accessed if required
- otherwise the **ALU result** is simply **buffered** for one clock cycle


◆ *Write-back*

- the results generated by the instruction are written back to the **register file**

4.2 5-stage pipeline ARM organization

■ Data forwarding

- ◆ Because instruction execution is spread across three pipeline stages, the only way to **resolve data dependencies** without stalling the pipeline is to introduce forwarding paths
- ◆ **Forwarding paths** allow results to be passed between stages as soon as they are available
- ◆ Consider the following code sequence:
LDR **rN**, [. .] ; load rN from somewhere
ADD r2, r1, **rN** ; and use it immediately
- ◆ The **only way to avoid this stall** is to encourage the **compiler not** to put a dependent instruction immediately after a load instruction



4.3 ARM instruction execution

- Data path organization: **Fig. 4.1** (3-stage pipeline)

- **Data processing instructions**

- ◆ A data processing instruction requires **two operands**, one of which is always a **register** and the other is either a **second register** or an **immediate value**
- ◆ The **second operand** is passed through the **barrel shifter** where it is subject to a general shift operation

ADD r3, r2, r1, LSL #3 ; r3 := r2 + 8 × r1

- ◆ The **result** from the ALU is written back into the destination register (the condition code register may be updated)
- ◆ All these operations take place in a **single clock cycle** as shown in **Fig. 4.5**
- ◆ Only the bottom eight bits (**bits [7:0]**) of the instruction are used in the **immediate value**





Fig. 4.5 Data processing instruction datapath activity



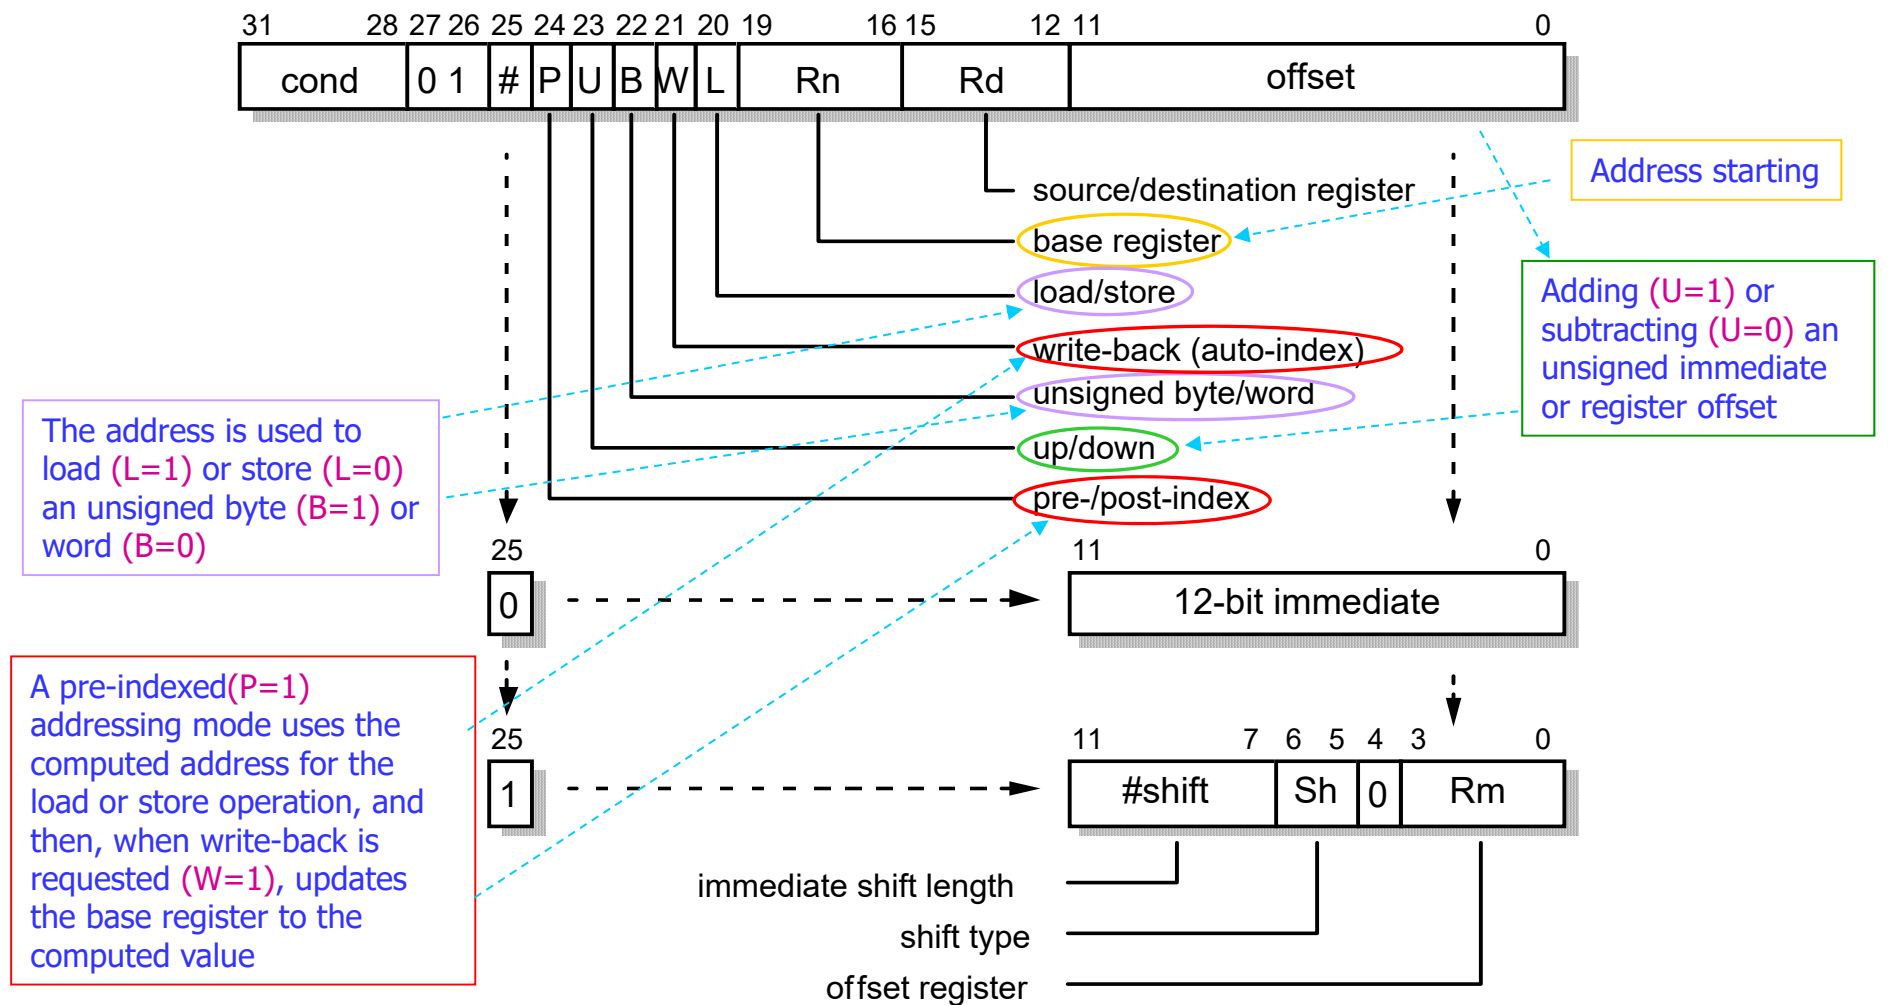
$$immediate = (0 \rightarrow 255) \times 2^{2n}, \quad \text{where } 0 \leq n \leq 12$$

4.3 ARM instruction execution

■ Data transfer instructions

- ◆ A data transfer (load or store) instruction computes a **memory address** in a manner very similar to the way a data processing instruction computes its result
- ◆ A **register** is used as the **base address**, to which is added (or from which is subtracted) an **offset** which again may be **another register** or an **immediate value** (12-bit)
- ◆ The address is sent to the **address register**, and in a **second cycle** the data transfer takes place (Fig. 4.3)
- ◆ Rather than leave the datapath largely idle during the data transfer cycle, the **ALU** holds the **address** components from the first cycle and is available to compute an **auto-indexing** modification to the base register if this is required

```
LDR    r0, [ r1 , #4 ]!      ; r0 := mem32[ r1 + 4 ]  
                                ; r1 := r1 + 4
```



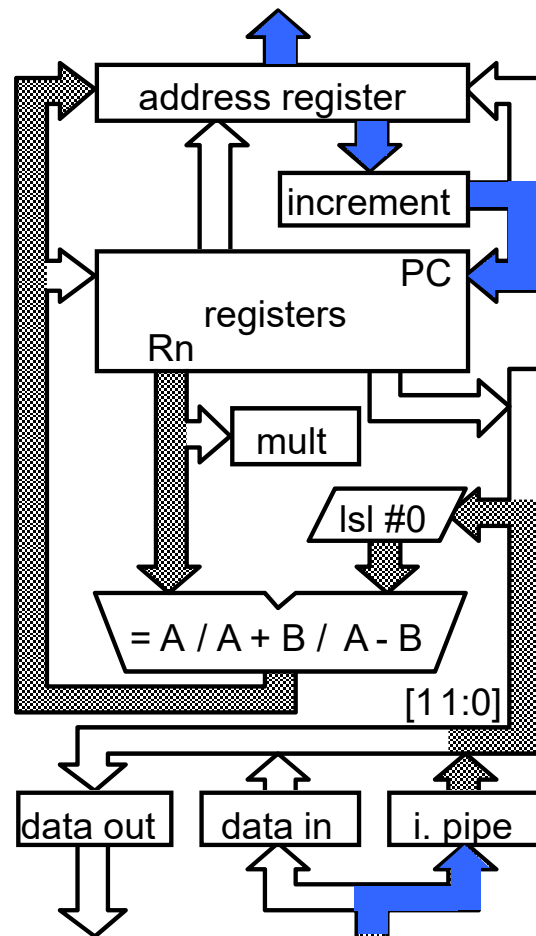

LDRLS pc, [r1, r0, LSL #2] Chap.3, p.35

Fig 5.9 Single word and unsigned byte data transfer instruction binary encoding

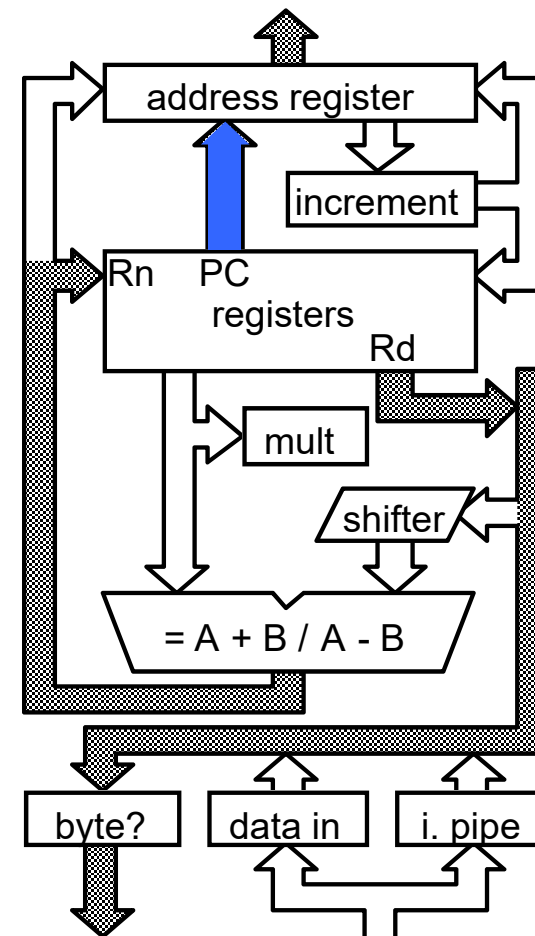
4.3 ARM instruction execution

■ Data transfer instructions (Cont.)

- ◆ The datapath operation for the **two cycles** of a data store instruction (**STR**) with an **immediate offset** are shown in **Fig. 4.6**
- ◆ The incremented PC value is stored in the register bank at the end of the **first cycle** so that the address register is free to accept the **data transfer address** for the second cycle
- ◆ At the end of **second cycle** the **PC** is fed back to the address register to allow instruction prefetching to continue
- ◆ The value sent to the address register in a cycle is the value used for the memory access in the following cycle
- ◆ The address register is a pipeline register between the processor datapath and the external memory



(a) 1st cycle - compute address



(b) 2nd cycle - store data & auto-index

Fig. 4.6 STR (store register) datapath activity

4.3 ARM instruction execution

■ Branch instructions

- ◆ Compute the **target address** in the **first cycle** as shown in Fig. 4.7
 - A 24-bit immediate field is extracted from the instruction and then **shifted left two** bit positions to give a word-aligned offset which is added to the PC
- ◆ The **second cycle**: the result is issued as an **instruction fetch address**, and while the instruction **pipeline refills** the **return address** is copied into the link register (r14) if this is required (instruction is a '**branch with link**')
- ◆ The **third cycle**, which is required to **complete the pipeline refilling**, is also used to make a small **correction to the value stored in the link register** in order that it points directly at the instruction which follows the branch
 - This is necessary because r15 contains **pc+8** whereas the address of the next instruction is **pc+4**

