

Chap. 5 The ARM Instruction Set

Data types

- ARM processors supports six data types:
 - 8-bit signed and unsigned bytes
 - 16-bit signed and unsigned half-words; these are aligned on 2-byte boundaries
 - 32-bit signed and unsigned words; these are aligned on 4-byte boundaries
- ARM instructions are all 32-bit words and must be word-aligned
- Thumb instructions are half-words and must be aligned on 2-byte boundaries
- All ARM operations are on 32-bit operands; the shorter data types are only supported by data transfer instructions

Memory organization

- There are two ways to store words in a byte-addressed memory, depending on whether the least significant byte is stored at a lower or higher address than the next most significant byte
- The issue causes significant practical difficulties when datasets are transferred between machines of opposite orderings
- Most ARM chips remain strictly neutral in the dispute and can be configured to work with either memory arrangement, though they default to little-endian
- This book we will assume a little-endian ordering

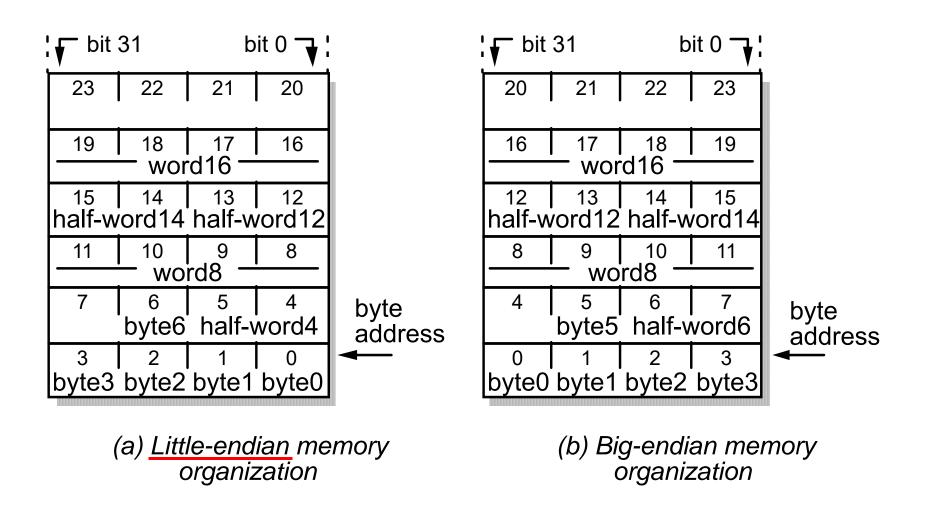


Fig. 5.1 Little- and big-endian memory organizations

Privileged modes

- The current operating mode is defined by the bottom five bits of the CPSR
- The relevant shaded registers shown in Fig. 2.1 replace the corresponding user registers and the current SPSR (Saved Program Status Register) also becomes accessible

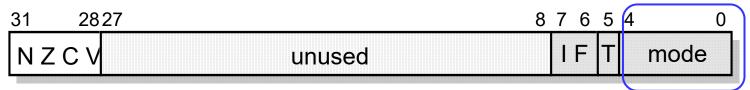


Fig. 2.2 ARM CPSR format

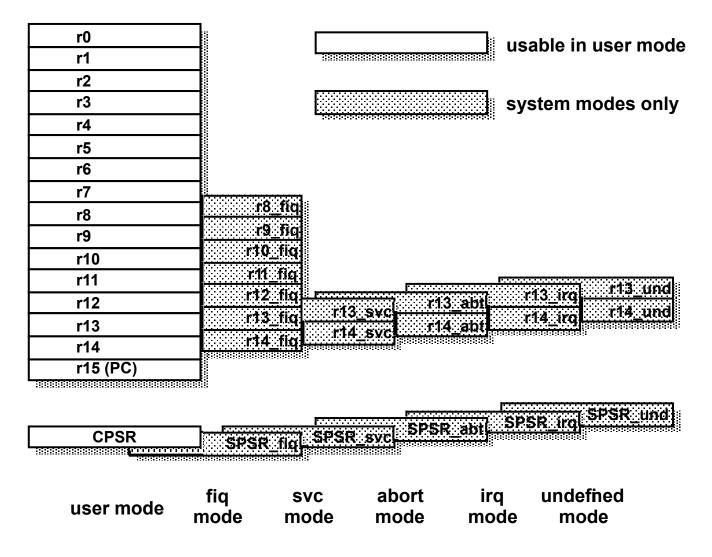


Fig. 2.1 ARM's visible registers

CPSR[4:0]	Mo de	Use	Registers
10000	User	Normal user code	user
10001	FIQ	Processing fast interrupts	_fiq
10010	IRQ	Processing standard interrupts	_irq
10011	SVC	Processing software interrupts (SWIs)	_svc
10111	Abort	Processing memory faults	_abt
11011	Undef	Handling undefined instruction traps	_und
11111	System	Running privileged operating system tasks	user

Table 5.1 ARM operating modes and register usage

The SPSRs

- This register is used to save the state of the CPSR (Current Program Status Register) when the
- mode is entered in order that the user state can be fully restored when the user process is resumed
- Often the SPSR may be untouched from the time the privileged mode is entered to the time it is used to restore the CPSR
- If the privileged software is to be re-entrant then the SPSR must be copied into a general register and saved

5.2 Exceptions

ARM exceptions

- Exceptions generated as the direct effect of executing an instruction – Software interrupts, undefined instructions and prefetch aborts (instructions that are invalid due to a memory fault occurring during fetch)
- Exceptions generated as a side-effect of an instruction data aborts (a memory fault during a load or store data access)
- Exceptions generated externally, unrelated to the instruction flow – Reset, IRQ and FIQ fall into this category

5.2 Exception Entry

- When an exception arises, ARM completes the current instruction as best it can and then departs from the current instruction sequence to handle the exception
- The processor performs the following sequence of actions:
 - It changes to the operating mode corresponding to the particular exception
 - It saves the address of the instruction following the exception entry instruction in r14 of the new mode
 - It saves the old value of the CPSR in the SPSR of the new mode
 - It disables IRQs by setting bit 7 of the CPSR
 - IRQs are disabled when any exception occurs.
 - FIQs are disabled when a FIQ occurs, and on reset
 - It forces the PC to begin executing at the relevant vector address (is normally a branch) in Table 5.2
 - normally the vector address will contain a branch to the relevant routine
- The two <u>banked registers</u> (r14 and r13) in each of the privileged modes are used to hold the return address and a stack pointer NSYSU CSE 組合語言與微處理機 Chap. 5

Exception	Mo de	Vector address
Reset	SVC	0x00000000
Undefined instruction	UND	0x00000004
Software interrupt (SWI)	SVC	0x00000008
Prefetch abort (instruction fetch memory fault)	Abort	0x000000C
Data abort (data access memory fault)	Abort	0x0000010
IRQ (normal interrupt)	IRQ	0x00000018
FIQ (fast interrupt)	FIQ	0x000001C

Table 5.2 Exception vector addresses

Startup code

- Startup code for C/C++ program consists of the following actions
 - Disable all interrupt
 - Copy any initialized data from ROM to RAM
 - Zero the uninitialized data area
 - Allocate space for and initialize the stack
 - Initialize the processor's stack pointer
 - Create and initialize the heap
 - Execute the constructors and initializers for all global variables (C++ only)
 - Enable interrupts
 - Call main

5.2 Exceptions

Exception return

- Any modified user registers must be restored from the handler's stack
- The CPSR must be restored from the appropriate SPSR
- The PC must be changed back to the relevant instruction address in the user instruction stream
- The last two steps cannot be carried out independently.
- The case where the return address is in r14 (S: destination is pc):
 - To return from a SWI or undefined instruction trap use

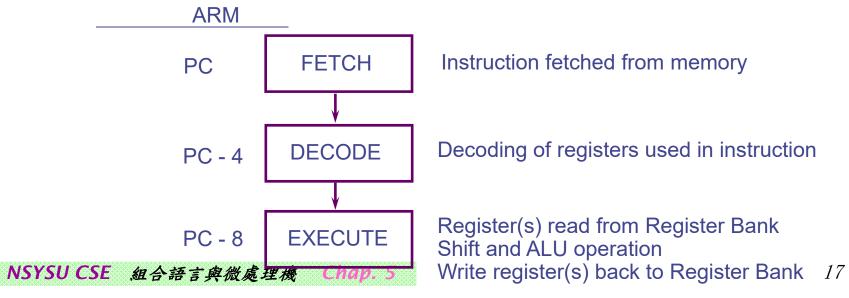
- To return from an IRQ, FIQ or prefetch abort use

- To return from a data abort to retry the data access use

- The return address out onto a stack:
 - LDMFD r13!, $\{r0 r3, pc\}^{\wedge}$; restore and return

The Instruction Pipeline

Exception	Address	Use	
Reset	-	Ir is not defined	
Data abort	Ir-8	Points to the instructions that caused data abort	
FIQ	Ir-4	Return to next instruction	
IRQ	Ir-4	Return to next instruction	
Prefetch Abort	Ir-4	Points to the instruction that cause prefetch aboint	
SWI	Ir	Return to next instruction after SWI	
Undefined Instruction	Ir	Return to next instruction after undefined instruction	



5.2 Exceptions

Exception priorities

- Since multiple exceptions can arise at the same time it is necessary to define a priority order to determine the order in which the exceptions are handled:
 - 1. Reset (highest priority)
 - 2. Data abort
 - 3. FIQ
 - 4. IRQ
 - 5. Prefetch abort
 - 6. SWI, undefined instruction

Some instruction execution example

Example1

- PRE
 - cpsr=nzcvqiFt_USER
 - r1 = 0x00000001
 - SUBS r1, r1, #1
- POST
 - cpsr=nZCvqiFt_USER
 - r1 = 0x00000000

Example 2

- PRE
 - cpsr=nzcvqiFt_USER
 - r0 = 0x00000000
 - r1 = 0x80000004
 - MOVS r0, r1, LSL #1
- POST
 - cpsr=nzCvqiFt_USER
 - r0 = 0x00000008
 - r1 = 0x80000004

Example3

- PRE
 - cpsr=nzcVqift USER
 - pc = 0x00008000
 - Ir = 0x003fffff; Ir = r14
 - r0 = 0x12
 - 0x00008000 SWI 0x123456
- POST
 - cpsr=nzcVqift_SVC
 - spsr=nzcVqift_USER
 - pc = 0x00000008
 - Ir = 0x00008004;
 - r0 = 0x12

5.3 Conditional execution

Conditional execution

 The condition field occupies the top four bits of the 32-bit instruction field:

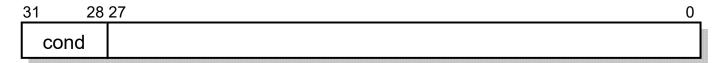


Fig. 5.2 The ARM condition code field

- Each of the 16 values of the condition field causes the instruction to be executed or skipped according to the values of the N, Z, C and V flags in the CPSR
- The conditions are given in Table 5.3
- The 'always' condition (AL) may be omitted since it is the default condition

Opcode [31:28]	Mnemonic extension	Interpretation	Status flag state for execution
0000	EQ	Equal / equals zero	Z set
0001	NE	Not equal	Z clear
0010	CS/HS	Carry set / unsigned higher or same	C set
0011	CC/LO	Carry clear / unsigned lower	C clear
0100	MI	Minus / negative	N set
0101	PL	Plus / positive or zero	N clear
0110	VS	Overflow	V set
0111	VC	No overflow	V clear
1000	HI	Unsigned higher	C set and Z clear
1001	LS	Unsigned lower or same	C clear or Z set
1010	GE	Signed greater than or equal	N equals V
1011	LT	Signed less than	N is not equal to V
1100	GT	Signed greater than	Z clear and N equals V
1101	LE	Signed less than or equal	Z set or N is not equal to V
1 110	AL	Always	any
1111	NV	Never (do not use!)	none

Table 5.3 ARM condition codes

Opcode	Interpretation
[31:28]	D 1 / 1
0000	Equal / equals zero
0001	Not equal
0010	Carry set / unsigned higher or same
0011	Carry clear / unsigned lower
0100	Minus / negative
0101	Plus / positive or zero
0110	Overflow
0111	No overflow
1000	Unsigned higher
1001	Unsigned lower or same
1010	Signed greater than or equal
1011	Signed less than
1100	Signed greater than
1101	Signed less than or equal
1110	Always
1111	Never (do not use!)

Table 5.3 ARM condition codes

5.3 Conditional execution

The 'never' condition

- The 'never' condition (NV) should not be used
 - There are plenty of other ways to write no-ops
- They may use this area of the instruction space for other purposes in the future

Alternative mnemonics

 This indicates that there is more than one way to interpret the condition field

```
CMP r0, #5 ; if (r0 != 5) {

ADDNE r1, r1, r0 ; r1 := r1 + r0 - r2

SUBNE r1, r1, r2 ; }
```

5.4 Branch and Branch with Link (B, BL)

Binary encoding

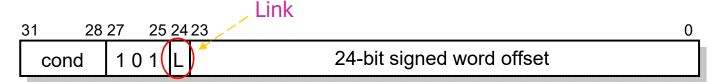


Fig. 5.3 Branch and Branch with Link binary encoding

Description

- Branch and branch with Link instructions cause the processor to begin executing instructions from an address computed by sign extending the 24-bit offset specified in the instruction
- The Branch with Link variant, which has the L bit (bit 24) set, also moves the address of the instruction following the branch into the link register (r14) of the current processor mode

5.4 Branch and Branch with Link (B, BL)

- Assembler format :
 - B{L} {<cond>} <target address>
 - 'L' specifies the branch and link variant
 - '<cond>' should be one of the mnemonic extensions
 - '<target address>' is normally a label in the assembler code

Example (conditional subroutine call)

```
CMP r0, #5; if r0 < 5
BLLT SUB1; then call SUB1
BLGE SUB2; else call SUB2
```

. .

5.5 Branch, Branch with Link and eXchange (BX, BLX)

BX, BLX

 These instructions are available on ARM chips which support the Thumb (16-bit) instruction set, and are a mechanism for switching the processor to execute Thumb instructions or for returning symmetrically to ARM and Thumb calling routines

Binary encoding

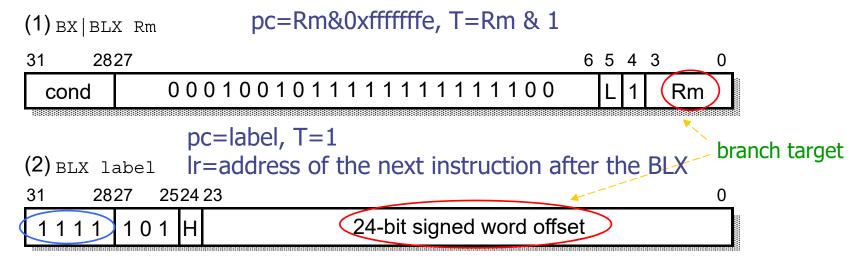
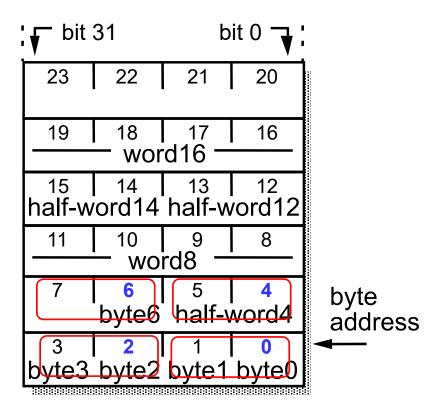


Fig. 5.4 Branch (with optional link) and exchange instruction binary encoding

5.5 Branch, Branch with Link and eXchange (BX, BLX)

- Description (Cont.)
 - In the *first format* the branch target is specified in a register, Rm
 - Bit[0] of Rm is copied into the T bit in the CPSR and bits[31:1] are moved into the PC:

- Bit 0 of offset If Rm[0] is 1, the processor switches to execute Thumb instructions and begins executing at the address in Rm aligned to a half-word boundary by clearing the bottom bit (....0)
 - If Rm[0] is 0, the processor continues executing ARM instructions and begins executing at the address in Rm aligned to a word boundary by clearing Rm[1] (....00)
 - In the second format the branch target is an address computed by sign extending the 24-bit offset specified in the instruction
 - Format (1) instructions may be executed conditionally or unconditionally, but format (2) instructions are executed unconditionally



Little-endian memory organization

5.5 Branch, Branch with Link and eXchange (BX, BLX)

Assembler format :

- 1: $B\{L\}X\{<cond>\}$ Rm
- 2: BLX <target address>

<target address> is normally a label

Example

```
CODE32 ; ARM code follows

BLX TSUB ; call Thumb subroutine

CODE16 ; start of Thumb code
; Thumb subroutine
; Thumb subroutine
; return to ARM code
```

Software interrupt

 The software interrupt instruction is used for calls to the operating system and is often called a 'supervisor call'

Binary encoding

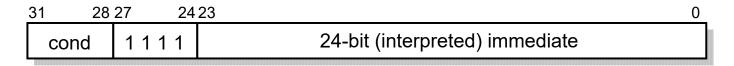


Fig. 5.5 Software interrupt binary encoding

Description

- The processor actions are:
 - 1. Save the address of the instruction after the SWI in r14_svc
 - 2. Save the CPSR in SPSR_svc
 - 3. Enter supervisor mode and disable IRQs by setting CPSR[4:0] to 10011₂ and CPSR[7] to 1
 - 4. Set the PC to 08₁₆ (is normally a branch to the SWI handler) and begin executing the instructions there
- To return to the instruction after the SWI the system routine must not only copy r14_svc back into the PC, but it must also restore the CPSR from SPSR_svc
- Assembler format

SWI{<cond>} <24-bit immediate>

Example

```
STROUT
                                   ; output following message
        BL
                "Hello World", &0a, &0d, 0
                                    : return to here
                r0, [r14], #1
                                   ; get character
STROUT
        LDRB
                                   ; check for end marker
        CMP
                r0, #0
                                   ; if not end, print ...
                SWI_WriteC
        SWINE
                                   ; .. and loop
        BNE
                STROUT
                r14, #3
                                   ; align to next word
        ADD
        BIC
                r14, #3
                pc, r14
        VOM
                                   ; return
```

Example

hexadecimal

```
Hellow, CODE, READONLY; declare code area
        AREA
SWI_WriteC
                EOU
                         60
                                     ; output character in r0
SWI_Exit
                EQU
                         &11
                                     ; finish program
        ENTRY
                                     ; code entry point
START
        ADR
                r1, TEXT
                                     ; r1 -> "Hello World"
LOOP
        LDRB
                r0, [r1], #1
                                     ; get the next byte
        CMP
                r0, #0
                                     ; check for text end
        SWINE
                SWI_WriteC
                                     ; if not end print ...
        BNE
                LOOP
                                     ; .. and loop back
        SWI
                SWI_Exit
                                     ; end of execution
TEXT
                "Hello World", &0a, &0d, 0
        END
                                     ; end of program source
```

SWI Exception

- The sample code of SWI handler
 - STMFD sp!, {r0-r12,r14}
 - LDR r10, [r14,#-4]
 - BIC r10, r10, #0xff000000
 - MOV r1,r13
 - MRS r2, spsr
 - STMFD r13!,{r2}
 - BL swi_jumptable
- The sample code of swi_jumptable
 - swi_jumptable
 - MOV r0, r10
 - B eventsSWIHandler

5.7 Data processing instructions

- ARM data processing instructions
 - are used to modify data values in registers
- Binary encoding: Fig. 5.6
- Description
 - When the instruction does not require all the available operands
 - MOV ignores Rn and CMP ignores Rd
 - The unused register field should be set to zero
 - Table 5.4

■ We can write: <u>MOV r1, -1</u> because the assembler will transfer it to a suitable instruction to replace the original MOV instruction.

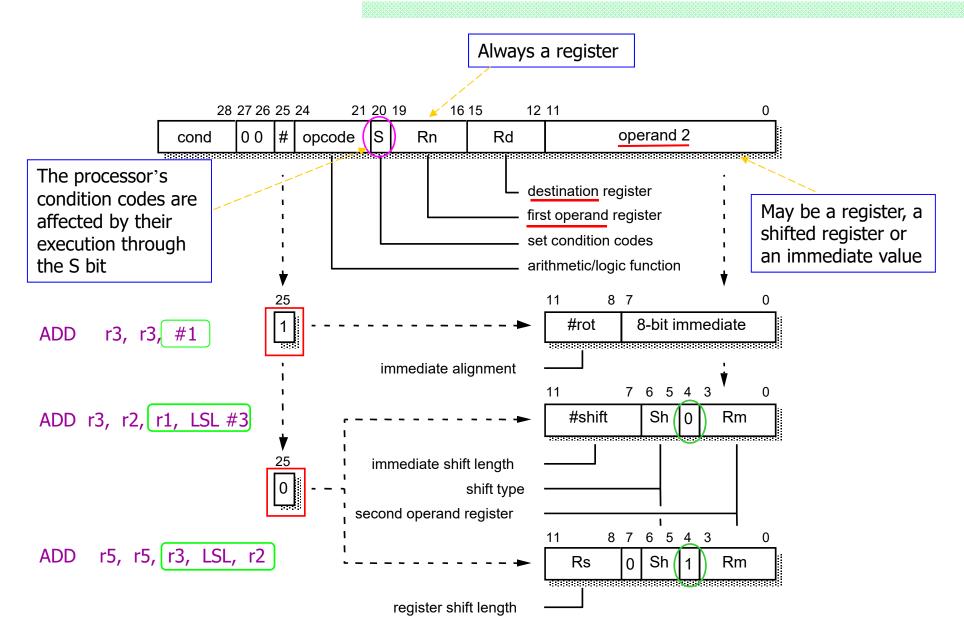
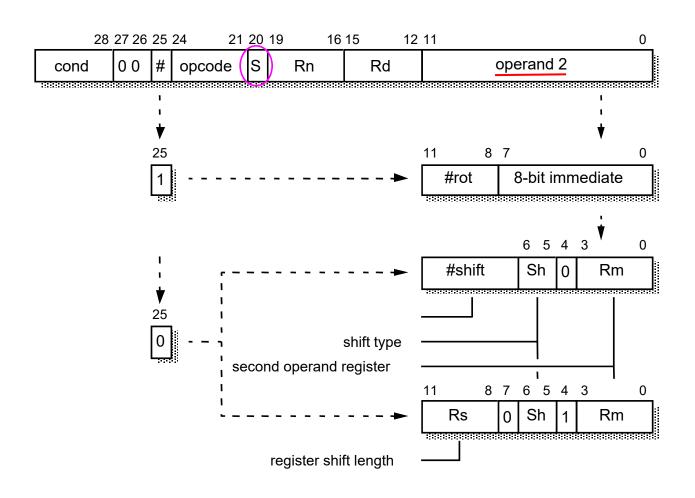


Fig. 5.6 Data processing instruction binary encoding



Op c o de [24:21]	Mnemonic	Meaning	Effect
0000	AND	Logical bit-wise AND	$Rd := Rn \ AND \ Op 2$
0001	EOR	Logical bit-wise exclusive OR	Rd := Rn EOR Op2
0010	SUB	Subtract	Rd := Rn - Op2
0011	RSB	Reverse subtract	Rd := Op2 - Rn
0100	ADD	Add	Rd := Rn + Op2
0101	ADC	Add with carry	Rd := Rn + Op2 + C
0110	SBC	Subtract with carry	Rd := Rn - Op2 + C - 1
0111	RSC	Reverse subtract with carry	Rd := Op2 - Rn + C - 1
1000	TST	Test	Scc on Rn AND Op2
1001	_TEQ	Test equivalence	Scc on Rn EOR Op2
1010	CMP	Compare	Scc on Rn - Op2
1011	CMN	Compare negated	Scc on Rn + Op2
1100	ORR	Logical bit-wise OR	Rd := Rn OR Op2
1101	MOV	Move	Rd := Op2
1110	BIC	Bit clear	Rd := Rn AND NOT Op 2
1111	MVN	Move negated	Rd := NOT Op2

Table 5.4 ARM data processing instructions

5.7 Data processing instructions

Assembler format

- <op>{<cond>}{S}Rd, Rn, #<32-bit immediate>
- <op>{<cond>}{S}Rd, Rn, Rm, {<shift>}
- The S bit controls the effect of the instruction on the CPSR (N, Z, C and V)

Examples

```
    SUBS r2, r2, #1
    BEQ LABEL ; branch if r2 zero
```

5.8 Multiply instructions

- Multiply instructions
 - The result is a 64-bit product
 - Some forms of the instruction
 - Store the full result into two independently specified registers
 - Store only the least significant 32 bits into a single register

Binary encoding

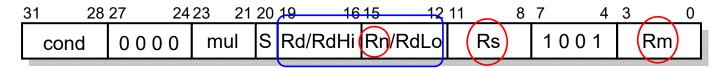


Fig. 5.7 Multiply instruction binary encoding

Description (Table 5.5)

5.8 Multiply instructions

Assembler formats

 Instructions that produce the least significant 32 bits MUL{<cond>}{S} Rd, Rm, Rs
 MLA{<cond>}{S} Rd, Rm, Rs, Rn

Produce the full 64-bit result
 <mul>{<cond>}{S} RdHi, RdLo, Rm, Rs

- Rd, RdHi and RdLo should be distinct from Rm, and RdHi and RdLo should not be the same register
- The 64-bit multipiles are available only on ARM7 versions with an 'M' in their name (ARM7DM, ARM7TM) and subsequent processors

31 28	27 24	23 21	20	19 16	15 12	11	\sim	8	7	4	3		0
cond	0000	mul	S	Rd/RdHi	Rn/RdLo	(Rs		100	1	(Rm	

Op c o de	Mnemonic	Meaning	Effect	
[23:21]	-11			
000	MUL	Multiply (32-bit result)	Rd := (Rm	n * Rs) [31:0]
001	MLA	Multiply-accumulate (32-b	it result) Rd := (Rm)	n * Rs + Rn) [31:0]
100	UMULL	Unsigned multiply long	RdHi:RdL	o := Rm * Rs
101	UMLAL	Unsigned multiply-accumu	ılate long / RdHi:RdL	o += Rm * Rs
110	SMULL	Signed multiply long	RdHi:RdL	.o:≡Rm * Rs
111	SMLAL	Signed multiply-accumula	te long RdHi:RdL	√ += Rm * Rs
				The state of the s
		//	/	
		Table 5.5 Multiply instr	uctions	Accumulation
	is the 64 -b	nit number	/	
	formed by	concatenating	Simple assignment	
	RdHi and R	.dLo	is denoted	

5.9 Count leading zeros (CLZ – architecture v5T only)

Binary encoding

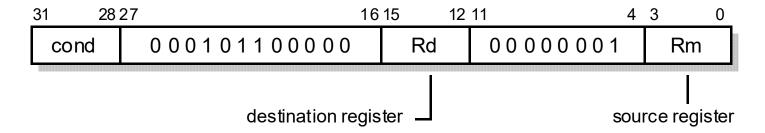


Fig. 5.8 Count leading zeros instruction binary encoding

Description

- The instruction sets Rd to the number of the bit position of the most significant 1 in Rm
- If Rm is zero Rd will be set to 32

Assembler format

CLZ{<cond>} Rd, Rm

5.9 Count leading zeros (CLZ – architecture v5T only)

Example

```
MOV r0, #&100
CLZ r1, r0 ; r1:=23
```

Notes

 Only processors that implement ARM architecture v5T support the CLZ instruction

5.10 Single word and unsigned byte data transfer instructions

- Description (Fig. 5.9)
- Assembler format
 - Pre-indexed form LDR|STR{<cond>}{B} Rd, [Rn, <offset>]{!}
 - Post-indexed form LDR|STR{<cond>}{B}{T} Rd, [Rn], <offset>
 - PC-relative LDR|STR{<cond>}{B} Rd, LABEL

Examples

```
    Pre-indexed: LDR r0, [r1, #4]!; r0 := mem<sub>32</sub>[r1 + 4]; r1 := r1 + 4
    Post-indexed: LDR r0, [r1], #4 ; r0 := mem<sub>32</sub>[r1]; r1 := r1 + 4
    PC-relative: LDR r1 LIAPTADD : range: 4 Kbytes
```

LDR r1, UARTADD ; range: 4 Kbytes

LDRB r0, [r1] ; r0 := mem_{8} [r1]

UARTADD & &1000000

5.10 Single word and unsigned byte data transfer instructions

- PC should not be used as the offset register
- Loading a byte into the PC should be avoided
- Storing the PC to memory be avoided if possible
- In general Rd, Rn and Rm should be distinct registers

```
LDR r0, [r1] ; r0 := mem_{32}[r1] P=0, W=0

LDR r0, [r1,#4] ; r0 := mem_{32}[r1 + 4] P=1, W=0 pre-indexed

LDR r0, [r1,#4]! ; r0 := mem_{32}[r1 + 4] P=1, W=1 pre-indexed

; r1 := r1 + 4 P=0, W=1

; r0 := mem_{32}[r1] P=0, W=1

; r1 := r1 + 4
```

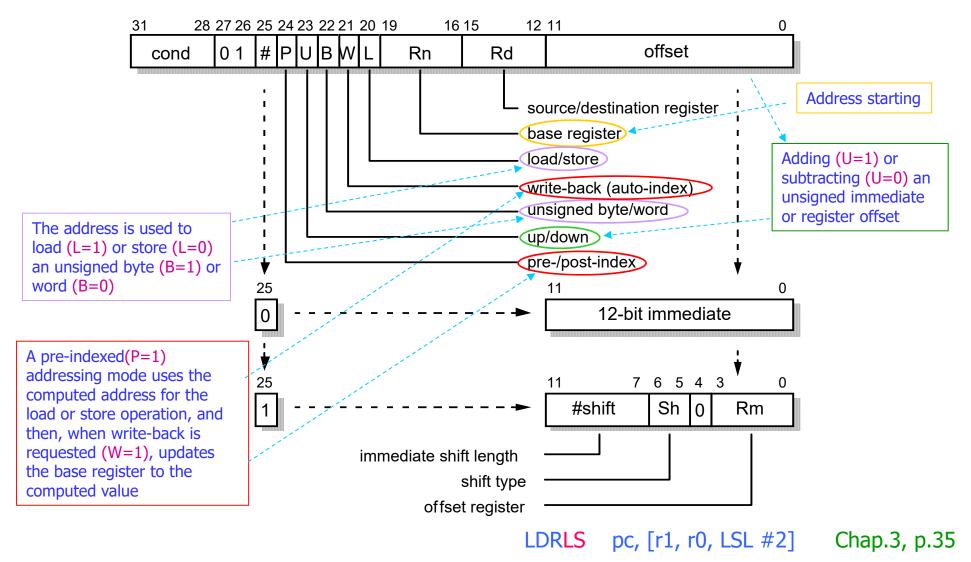


Fig 5.9 Single word and unsigned byte data transfer instruction binary encoding

5.11 Half-word and signed byte data transfer instructions

- Binary encoding (Fig. 5.10)
- Description
 - The only relevant forms of this instruction format are
 - Load signed byte, signed half-word or unsigned half-word
 - Store half-word
 - An unsigned value is zero-extended to 32 bits when loaded
 - A signed value is extended to 32 bits by replicating the most significant bit of the data
- Assembler format signed half-word
 - The pre-indexed form
 LDR|STR{<cond>}H|SH|SB Rd, [Rn, <offset>]{!}
 - The post-indexed form
 LDR|STR{<cond>}H|SH|SB Rd, [Rn], <offset>

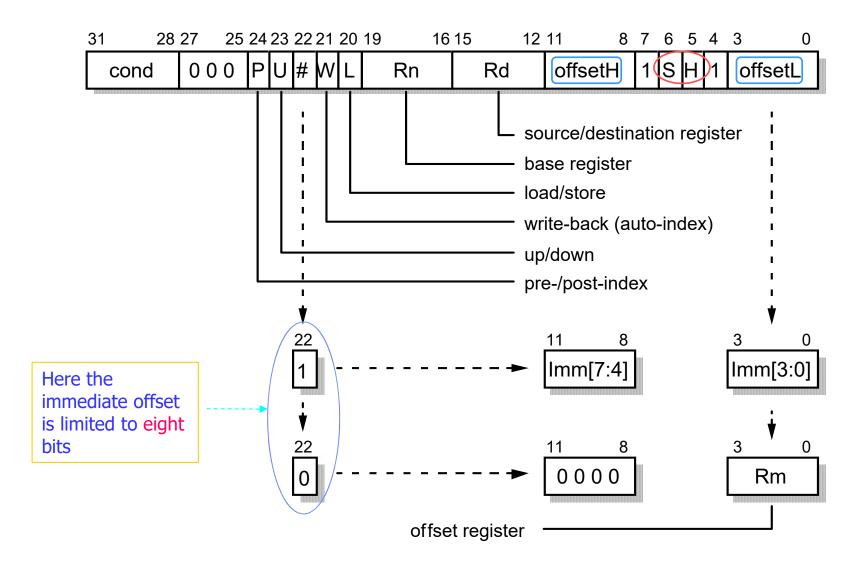


Fig. 5.10 Half-word and signed byte data transfer instruction binary encoding

5.11 Half-word and signed byte data transfer instructions

Notes

All half-word transfers should use half-word aligned address

<u>S</u>	H	Data type
1	0	Signed byte
0	1	Unsigned half-word
1	1	Signed half-word

Table 5.6 Data type encoding

Example

■ To expand an array of signed half-words into an array of words

ADR r1, ARRAY1

ADR r2, ARRAY2

ADR r3, ENDARR1

LOOP LDRSH r0, [r1], #2

• STR r0, [r1], #4

• CMP r1, r3

BLT LOOP

5.12 Multiple register transfer instructions

Binary encoding (Fig. 5.11)

Description

 The register list in the bottom 16 bits of the instruction includes a bit for each visible register, with bit 0 controlling whether or not r0 is transferred....

```
LDMIA r0!, {r2-r9} ; IA: increment after W=1

STMFD r13!, {r2-r9} ; FD: full descending

LDMIA r1, {r0, r2, r5} ; W=0
```

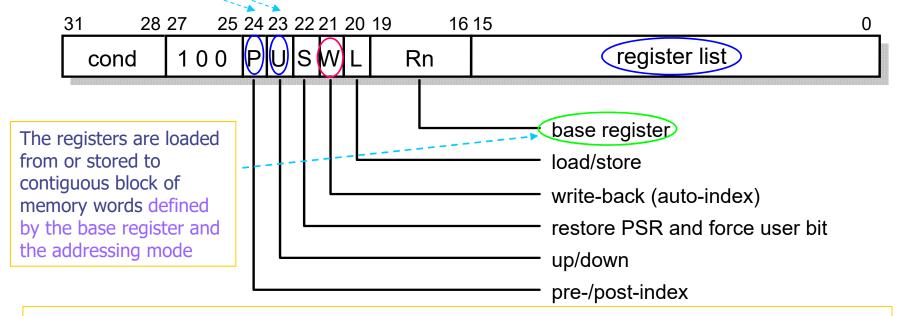
Normal form

LDM|STM{<cond>}<add mode> Rn{!}, <registers>

- In a non-user mode, the CPSR may be restored
 LDM{<cond>}<add mode> Rn{!}, <registers + PC>^
- In a non-user mode, the user registers may be saved or restored LDM|STM{<cond>}<add mode> Rn, <registers - PC>^

```
Incremented (U=1) or decremented (U=0)
Before (P=1) or after (P=0) each transfer
```

```
LDMIA r0!, \{r2-r9\} ; IA: increment after U=1, P=0, W=1 STMFD r13!, \{r2-r9\} ; FD: full descending (decrement before) U=0, P=1, W=1 LDMIA r1, \{r0, r2, r5\} ; W=0
```



Special forms:

If the PC is in the register list of a load multiple and the S bit is set, the SPSR of the current mode will be copied into the CPSR, giving an atomic return and restore state instruction.

```
SUB1 STMFD r13!, {r0-r2,r14}; save work regs & link

BL SUB2

...

LDMFD r13!, {r0-r2,pc}; restore work regs & return
```

5.12 Multiple register transfer instructions

- Specifying the PC in an STM should be avoided
- The base register may be specified in the transfer list, but writeback should not be specified in the same instruction

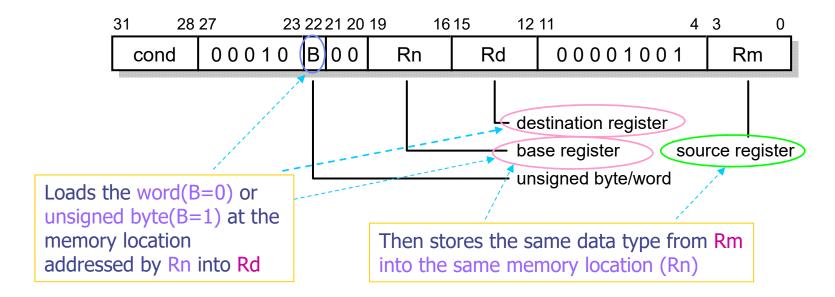
		Ascending		Desce	nding
		Full	Empty	Full	Empty
	Before	STMIB			LDMIB
Increment		STMFA			LDMED
'	After		STMIA	LDMIA	
			STMEA	LDMFD	
	Before		LDMDB	STMDB	
Decrement			LDMEA	STMFD	
'	After	LDMDA			STMDA
		LDMFA			STMED

Table 3.1

5.13 Swap memory and register instructions (SWP)

- Swap instructions
 - combine a load and a store of a word or an unsigned byte in a single instruction
 Mem[Rn]=Rm, Rd=Mem[Rn]
- Binary encoding

SWP{<cond>}{B} Rd, Rm, [Rn] SWPB r1, r1, [r0]; exchange byte



Rd and Rm may be the same register

5.13 Swap memory and register instructions (SWP)

Assembler format

```
SWP{<cond>}{B} Rd, Rm, [Rn]
```

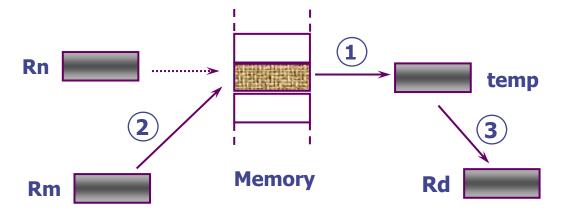
Example

```
ADR r0, SEMAPHORE SWPB r1, r1, [r0] ; exchange byte
```

- The PC should not be used as any of the registers in this instruction
- The base register (Rn) should not be the same as either the source (Rm) or the destination (Rd) register

Swap and Swap Byte Instructions

- Atomic operation of a memory read followed by a memory write which moves byte or word quantities between registers and memory.
- Syntax:
 - SWP {<cond>} {B} Rd, Rm, [Rn]



- Thus to implement an actual swap of contents make Rd = Rm.
- The compiler cannot produce this instruction.

SWAP instruction

The swap instruction is a special case of a load-store instructions. This instruction is an atomic operation.

PRE

- mem32[0x9000] = 0x12345678
- r0 = 0x00000000
- r1 = 0x11112222
- r2 = 0x00009000
- SWP r0, r1, [r2]

POST

- mem32[0x9000] = 0x11112222
- r0 = 0x12345678
- r1 = 0x11112222
- r2 = 0x00009000

Example of SWAP

Semaphore:

Provide a mutual exclusive access control to the shared resource.

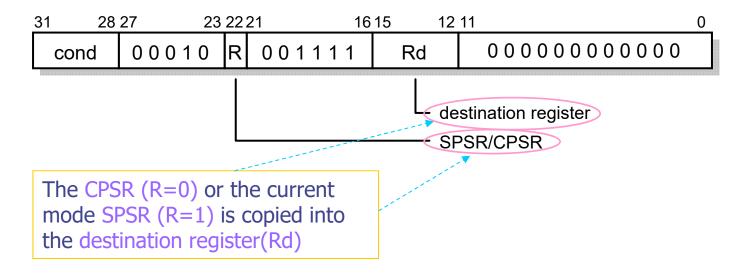
spin

- MOV r1, = semaphore
- MOV r2, #1
- SWP r3, r2, [r1]
- CMP r3, #1
- BEQ spin

5.14 Status register to general register transfer instructions

- When it is necessary to save or modify the contents of the CPSR or the SPSR of the current mode
 - Those contents must first be transferred into a general register
 - The selected bits modified
 - The value returned to the status register

Binary encoding



5.14 Status register to general register transfer instructions

Assembler format

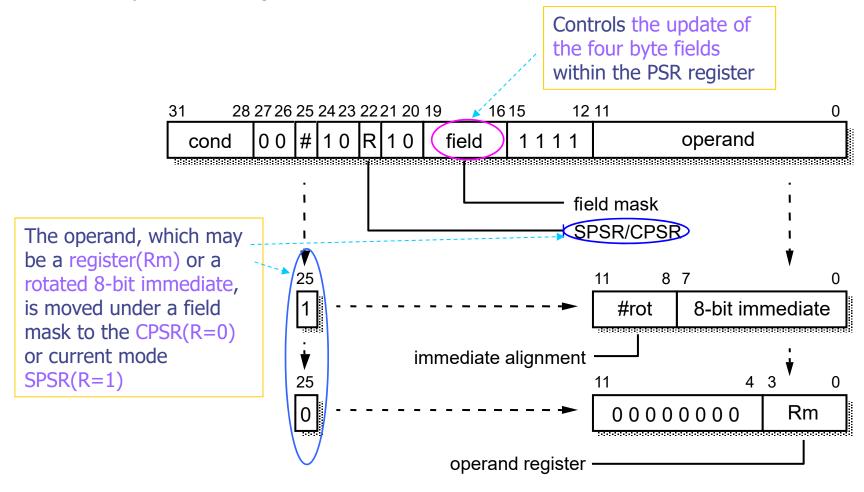
MRS{<cond>} Rd, CPSR|SPSR

Notes

 The SPSR form should not be used in user or system mode since there is no accessible SPSR in those modes

5.15 General register to status register transfer instructions

Binary encoding



5.15 General register to status register transfer instructions

Assembler format

```
MSR{<cond>} CPSR_f|SPSR_f, #<32-bit immediate>
MSR{<cond>} CPSR_<field>|SPSR_<field>, Rm

<field> is one of:

- c PSR[7:0] - x PSR[15:8]

- s PSR[23:16] - f PSR[31:24]
```

Examples

MRS r0, CPSR ; move the CPSR to r0 ORR r0, r0, #&20000000 ; set bit 29 of r0 MSR CPSR_f, r0 ; move back to CPSR

- When an immediate operand is used, only PSR[31:24] may be selected.
- Attempts to modify any of CPSR[23:0] whilst in user mode have no effect

31 28	27 8	7	6	5	4	0
NZCV	unused	ı	F	Т	mode	

Program status register instructions

■ The following example shows how to enable IRQ interrupts by clearing the I masks in the c field of CPSR.

PRE

- cpsr = nzcvqIFt_SVC
- MRS r1, cpsr
- ◆ BIC r1, r1, #0x80
- MSR cpsr_c, r1

POST

cpsr = nzcvqiFt_SVC

Why can user-level code not disable interrupts

The following code illustrates how a malicious user destroies all the currently active programs.

- The OS cannot regain control.
- The OS usually establishes a regular periodic interrupt from a hardware timer.

5.16 Coprocessor instructions

Coprocessor instructions

 The most common use of a coprocessor is the system coprocessor used to control on-chip functions such as the cache and memory management unit on the ARM720

Coprocessor registers

 ARM coprocessors have their own private register sets and their state is controlled by instructions that mirror the instructions that control ARM registers

Coprocessor data operations

 Coprocessor data operations are completely internal to the coprocessor and cause a stage change in the coprocessor registers

Coprocessor data transfers

- The ARM generates the memory address, but the coprocessor controls the number of word transferred
- A coprocessor may perform some type conversion as part of the transfer

Coprocessor register transfers

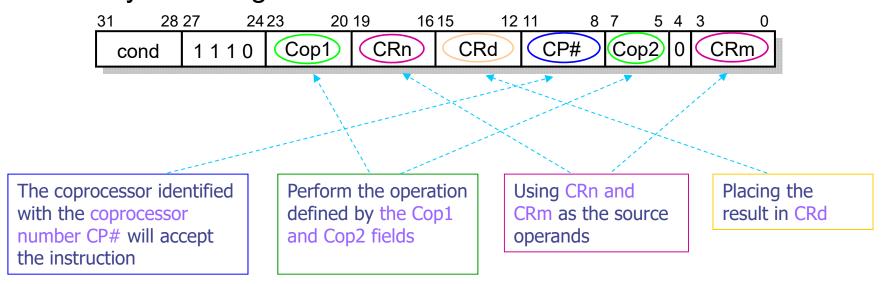
Move values between ARM and coprocessor registers

5.17 Coprocessor data operations

Coprocessor data operations

- The standard format follows the 3-address form of ARM's integer data processing instructions, but other interpretations of all the coprocessor fields are possible.
- Undefined instruction exception is raised if the instruction is not accepted.

Binary encoding



5.17 Coprocessor data operations

Assembler format

CDP{<cond>} <CP#>, <Cop1>, CRd, CRn, CRm{, <Cop2>}

Notes

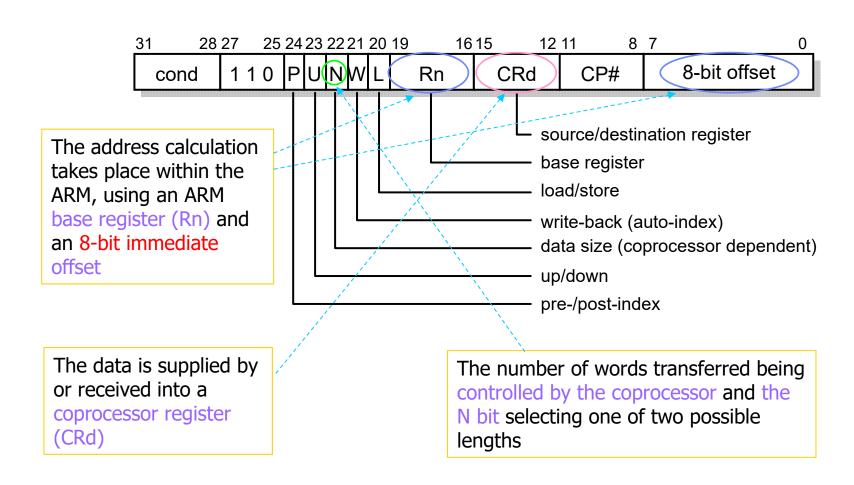
 The interpretation of the Cop1, CRn, CRd, Cop2 and CRm fields is coprocessor-dependent. It provides great flexibility.

Examples

- CDP p2, 3, C0, C1, C2
- CDPEQ p3, 6, C1, C5, C7, 4

5.18 Coprocessor data transfers

Binary encoding



5.18 Coprocessor data transfers

Assembler format

The pre-indexed formLDC|STC{<cond>}{L} <#CP>, CRd, [Rn, <offset>]{!}

The post-indexed form
 LDC|STC{<cond>}{L} <#CP>, CRd, [Rn], <offset>

Examples

- LDC p6, C0, [r1]
- STCEQL p5, C1, [r0], #4

- If the address is not word-aligned the two least significant bits will be ignored
- During the data transfer the ARM will not respond to interrupt requests, so coprocessor designers should be careful not to compromise the system interrupt response time by allowing very long data transfer
- If no coprocessor accepts the instruction, the undefined instruction traps.

5.19 Coprocessor register transfers

- Coprocessor register transfers
 - Allow an integer generated in a coprocessor to be transferred directly into a ARM register or the ARM condition code flags
 - A floating-point FIX operation which returns the integer to an ARM register
 - A floating-point comparison which returns the result of the comparison directly to the ARM condition code flags
 - A FLOAT operation which takes an integer value from an ARM register and sends it to the processor
- Binary encoding (Fig. 5.17)

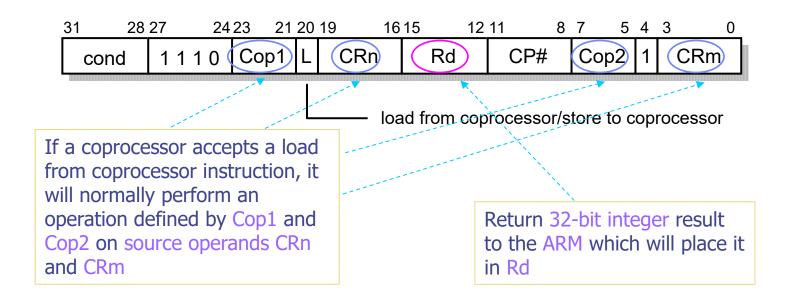


Fig. 5.17 Coprocessor register transfer instruction binary encoding

5.19 Coprocessor register transfers

Assembler format

- Move to ARM register from coprocessor
 MRC{<cond>} <CP#>, <Cop1>, Rd, CRn, CRm{, <Cop2>}
- Move to coprocessor from ARM register
 MCR{<cond>} <CP#>, <Cop1>, Rd, CRn, CRm{, <Cop2>}

Examples

- MCR p14, 3, r0, C1, C2
- MRCCS p2, 4, r3, C3, C4, 6

- The coprocessor must perform some internal work to prepare a 32-bit value for transfer to the ARM, it will often be necessary for the coprocessor handshake to 'busy-wait' while the data is prepared
- Transfers from the ARM to the coprocessor are generally simpler since any data conversion work can take place in the coprocessor after the transfer has completed

5.20 Breakpoint instruction (BKPT)

- Breakpoint instruction (BKPT)
 - Breakpoint instructions are used for software debugging purposes
 - This instruction causes the processor to take a prefetch abort when the debug hardware unit is configured appropriately
- Binary description

31 28	27 20	19 16	15 12	11 8	7 4	3 0
1110	00010010	xxxx	xxxx	xxxx	0111	xxxx

- Only processors that implement ARM architecture v5T support the BRK instruction
- BRK instruction are unconditional

5.21 Unused instruction space

- Unused instruction space
 - not all of the 2³² instruction bit encodings have been assigned meanings
 - are available for future instruction set extensions
- Unused arithmetic instructions This would be a likely encoding, for example, for an integer divide instruction

31 28	27 22	21 20	19 16	15 12	11 8	7 4	3 0
cond	000001	ор	Rn	Rd	Rs	1001	Rm

Unused control instructions - The gaps here could be used to encode other instructions that affect the processor operating mode

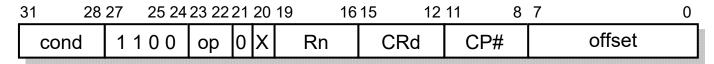
31 28	27 23	2221	20	19 16	15 12	11 8	7	6	4	3 0
cond	00010	op1	0	Rn	Rd	Rs		pp2	0	Rm
cond	00010	op1	0	Rn	Rd	Rs	0	op2	1	Rm
cond	00110	op1	0	Rn	Rd	#rot		3-bit	imı	mediate

5.21 Unused instruction space

Unused load/store instructions – These are likely to be used to support additional data transfer instructions, should these be required in the future

31	28 27	25	24	23	22	21	20	19 1	615	12	11	8	7	6	5	4	3		0
cond	0	0 0	Ρ	U	В	W	L	Rn		Rd	R	S	1	ор	1	1		Rm	

Unused coprocessor instructions – is likely to be used to support any additional coprocessor instructions



Undefined instruction space

5.22 Memory faults

- ARM processors allow the memory system to fault on any memory access
 - The memory system returns a signal that indicates that memory access has failed to complete correctly
 - The processor will then enter an exception handler and the system software will attempt to recover from the problem
- The most common sources of a memory fault in a generalpurpose machine are
 - Page absent: the addressed memory location has been paged out to disk
 - Page protected: the addressed memory location is temporarily inaccessible
 - Soft memory errors: a soft error has been detected in the memory
 - memory has a hardware error detector but relies on software error correction

5.22 Memory faults

Embedded systems

- A hard disk is usually unavailable
- The memory system is usually small
- Many embedded systems will not use memory faults at all

Memory faults

 The ARM handles memory faults detected during instruction fetches (prefetch aborts) and those detected during data transfers (data aborts) separately

Exception	Mo de	Vector address
Reset	SVC	0x00000000
Undefined instruction	UND	0x00000004
Software interrupt (SWI)	SVC	0x00000008
Prefetch abort (instruction fetch memory fault)	Abort	0x0000000C
Data abort (data access memory fault)	Abort	0x00000010
IRQ (normal interrupt)	IRQ	0x00000018
FIQ (fast interrupt)	FIQ	0x000001C

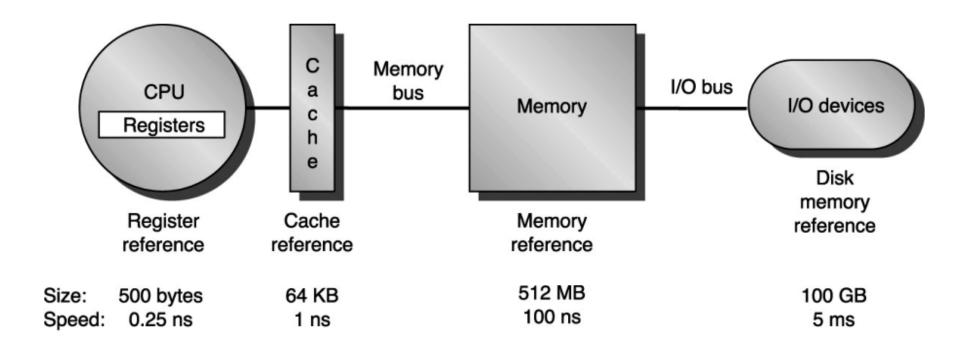


Fig. The levels in a typical memory hierarchy in embedded, desktop, and server computers

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	<1 KB	< 16 MB	< 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5000-10,000	1000-5000	20-150
Managed by	compiler	hardware	operating system	operating system/operator
Backed by	cache	main memory	disk	CD or tape

Fig. The typical levels in the hierarchy slow down and get larger as we move away from the CPU for a large workstation or small server

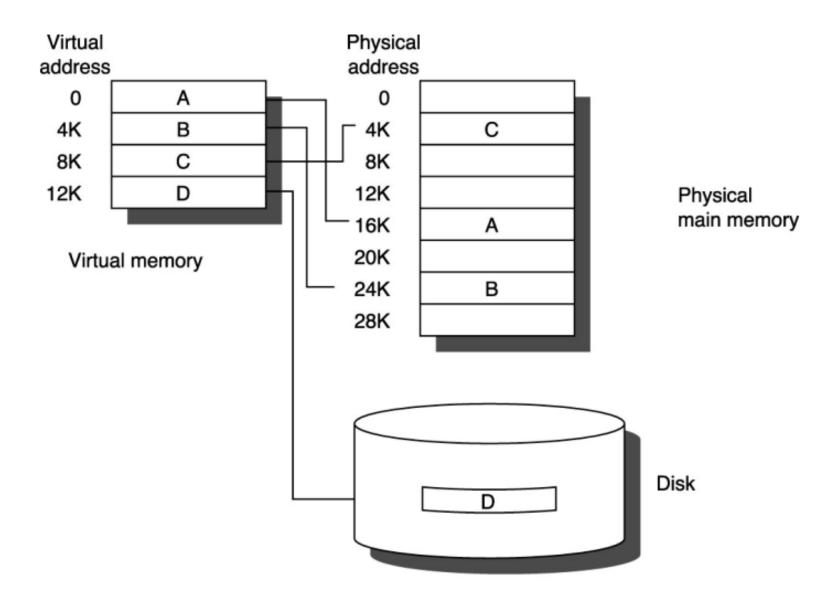


Fig. The logical program in its contiguous virtual address space in shown on the left.